

OPTIMIZATION FOR DEEP LEARNING

TOULOUSE SCHOOL OF ECONOMICS

Ludovic De Matteis

COURSE OVERVIEW

PLANNING


	Day	Time	Subject
1	13/10/2025	9:30 - 12:30	Basic definitions, Gradient Descent and Newton's method
2	27/10/2025	9:30 - 12:30	Practical - Gradient Descent and Newton's method
3	03/11/2025	9:30 - 12:30	Neural networks and stochastic gradient descent
4	10/11/2025	9:30 - 12:30	Practical - Neural Networks and digit recognition
5	17/11/2025	9:30 - 12:30	Alternative Neural Structures
6	24/11/2025	9:30 - 12:30	Practical - Alternative Neural structure, Adversial networks

LECTURE 2 -

OPTIMIZING NEURAL NETWORKS

SUMMARY

1. From perceptrons to multi-layers perceptron
 - Definition of a perceptron
 - Limits of single layer perceptrons
 - Extension to multi-layer perceptrons
2. Architecture and forward propagation
3. Backpropagation and derivatives
 - Backpropagation algorithm
 - Vanishing and exploding gradients
4. Optimization algorithms
 - Stochastic gradient descent
 - Mini batch gradient descent
 - Momentum and Adam optimizer



FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

MOTIVATIONS

- Neural networks are inspired by the structure of the human brain where neurons are interconnected to process information and learning occurs by adjusting the strength of these connections
- The perceptron is a fundamental building block of neural networks, representing a simplified model of a biological neuron
- In machine learning, we look for a parametrized model capable of approximating a function $f : \mathbb{R}^d \mapsto \mathbb{R}^n$ from examples

$$f_{\theta} \approx (x_i)y_i$$

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

HISTORY

- **1943:** McCulloch and Pitts propose a mathematical model of a neuron based on logical gates
- **1958:** Frank Rosenblatt invents the perceptron, an early neural network model capable of binary classification
- **1969:** Minsky and Papert publish “Perceptrons”, highlighting limitations of single-layer perceptrons, leading to a decline in neural network research
- **1986:** Rumelhart, Hinton, and Williams popularize the backpropagation algorithm, enabling training of multi-layer neural networks

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

DEFINITION OF A PERCEPTRON

- A perceptron is a simple computational unit that takes multiple input signals, applies weights to them, sums them up, and passes the result through an activation function to produce an output
- Mathematically, a perceptron can be represented as:

$$\text{output} = \Phi\left(\sum(\text{weights} * \text{inputs}) + \text{bias}\right)$$

or

$$y = \Phi(w^T x + b)$$


where:

- x are the **input** features
- w are the **weights**, the parameters learned during training
- b is the bias, an additional parameter to shift the activation function
- Φ is the **activation function** introducing non-linearity (e.g., step function, sigmoid, ReLU)

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

DEFINITION OF A PERCEPTRON

- The classical representation of a perceptron is as follows:



FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

LINEAR CLASSIFIER

- If we use a step activation function, the perceptron can be used as a linear classifier

$$\Phi(z) = 1 \text{ if } z \geq 0 \text{ else } 0$$

- The decision boundary is defined by the equation:

$$w^T x + b = 0$$

- In the first practical session, we implemented the perceptron using a sigmoid activation function (closer to the step function)

$$\Phi(z) = \frac{1}{1 + e^{-z}}$$

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

OPTIMIZING A PERCEPTRON

- We look for the optimal values of w and b
- We can use gradient descent to minimize a loss function, such as the mean squared error or cross-entropy loss (if everything is differentiable)
- Or we can use the following algorithm:
 1. Compute the output of the perceptron for each training example

$$\hat{y}_i = \Phi(w^T x_i + b)$$

2. Update the weights and bias based on the prediction error

$$w := w + \eta(y_i - \hat{y}_i)x_i$$

$$b := b + \eta(y_i - \hat{y}_i)$$

where η is the learning rate, controlling the step size of the updates

FROM PERCEPTRONS TO MULTI-LAYER PERCEPTRONS

LIMITS OF SINGLE-LAYER PERCEPTRONS

- Single-layer perceptrons can only learn linearly separable functions
- They cannot solve problems that require non-linear decision boundaries, such as the XOR problem
- To overcome these limitations, we extend the perceptron to multi-layer architectures, allowing for more complex representations and decision boundaries

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

MOTIVATIONS

- To obtain non linear functions, we combine multiple linear functions with non-linear activation functions
- The output will therefore be defined for a composition of L functions as:

$$y = \Phi_L(W_L \cdot \Phi_{L-1}(W_{L-1} \dots \Phi_1(W_1 x + b_1) \dots + b_{L-1}) + b_L)$$

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

ARCHITECTURE OF A MULTI-LAYER PERCEPTRON

- A multi-layer perceptron (MLP) consists of multiple layers of interconnected perceptrons (neurons)
- It typically includes an input layer, one or more hidden layers, and an output layer
- Each layer applies a linear transformation followed by a non-linear activation function (that can be different for each layer)
- The architecture allows the network to learn complex patterns and representations from the data
- The function computed by each layer l can be expressed as:


$$a^l = \Phi_l(W_l a^{l-1} + b_l)$$

with $a^0 = x$

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

ARCHITECTURE OF A MULTI-LAYER PERCEPTRON

- Example of a simple MLP with one hidden layer (we often omit the bias nodes in diagrams for clarity):



MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

FORWARD PROPAGATION

- Forward propagation is the process of computing the output of the MLP given an input
- It involves passing the input through each layer, applying the linear transformation and activation function at each step
- The steps for forward propagation are as follows:
 1. Initialize the input layer with the input features x
 2. For each layer l from 1 to L :
 - Compute the linear transformation:

$$z^l = W_l a^{l-1} + b_l$$

- Apply the activation function:

$$a^l = \Phi_l(z^l)$$

3. The final output is obtained from the output layer:

$$\hat{y} = a^L$$

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

FORWARD PROPAGATION ALGORITHM

- It can be simply implemented in Python as follows:

```
def forward_propagation(X, weights, biases):  
    a = X  
    for w, b, Phi in zip(weights, biases):  
        z = w @ a + b  
        a = phi(z)  
    return a
```

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

XOR PROBLEM

- The XOR problem is a classic example that illustrates the limitations of single-layer perceptrons and the capabilities of multi-layer perceptrons
- It can be solved using a MLP with a single hidden layer of two neurons
- It is possible to write by hand the formulation of the output

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

ACTIVATION FUNCTIONS

- The choice of activation function can significantly impact the performance and convergence of the network during training
- Common activation functions used in MLPs include:
 - Sigmoid:

$$\Phi(z) = \frac{1}{1 + e^{-z}}$$

Bounded, differentiable, but can suffer from vanishing gradients

- ReLU (Rectified Linear Unit):

$$\Phi(z) = \max(0, z)$$

Simple and efficient, helps mitigate vanishing gradients, but can suffer from dying ReLU problem

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

ACTIVATION FUNCTIONS

- Tanh:

$$\Phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Bounded between -1 and 1 , zero-centered, but can also suffer from vanishing gradients

- Leaky ReLU:

$$\Phi(z) = \max(\alpha z, z)$$

Similar to ReLU but allows a small gradient when the unit is not active

- Softmax:

$$\Phi(z_i) = \frac{e^{z_i}}{\sum(e^{z_j})}$$

Used in the output layer for multi-class classification problems

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

LOSS FUNCTIONS

- The loss function quantifies the difference between the predicted output and the true output
- Common loss functions used in training MLPs include:
 - Mean Squared Error (MSE):

$$L(y, \hat{y}) = \frac{1}{2} \sum ((y_i - \hat{y}_i)^2)$$

Used for regression tasks

- Binary Cross-Entropy Loss:

$$L(y, \hat{y}) = - \sum (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

Used for binary classification tasks

- Cross-Entropy Loss:

$$L(y, \hat{y}) = - \sum (y_i \log(\hat{y}_i))$$

Used for classification tasks, especially with softmax outputs

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

FEATURE DETECTION IN LAYERS

- In MLPs, each layer can be thought of as learning to detect different features of the input data
- Early layers may learn to detect simple features (e.g., edges in images), while deeper layers can learn more complex patterns (e.g., shapes, objects)
- This hierarchical feature learning allows MLPs to effectively model complex relationships in the data

MULTI-LAYER PERCEPTRON AND FORWARD PROPAGATION

EXERCISE

- Given

$$x = [0.5, -1.0]^T, \quad W^{(1)} = \begin{pmatrix} 1 & 2 \\ -1 & 0.5 \end{pmatrix}, \quad b^{(1)} = [0, 1]^T, \quad W^{(2)} = [1, -2], \quad b^{(2)} = [0]$$

- Draw the architecture of the MLP
- Compute $z^{(1)}, a^{(1)}, z^{(2)}, a^{(2)}$
- Compute the output of the MLP using ReLU activation function

BACKPROPAGATION AND DERIVATIVES

BACKPROPAGATION AND DERIVATIVES

MOTIVATIONS

- The goal of the learning process is to minimize the loss function $L(y, \hat{y})$ by adjusting the parameters of the network $\theta = \{W^{(l)}, b^{(l)}\}$

$$\min_{\theta} L(y, \hat{y}(\theta))$$

- To achieve this, we need to compute the gradients of the loss function with respect to each parameter in the network

$$\nabla_{\theta} L = \frac{\partial L}{\partial \theta}$$

- The backpropagation algorithm efficiently computes these gradients for Multi-Layer Perceptrons

BACKPROPAGATION AND DERIVATIVES

CHAIN RULE

- A neural network performs function composition

$$f(x) = f_L(f_{L-1}(\dots f_1(x)\dots))$$

- To compute the derivative of L with respect to a layer, we use the chain rule:

$$\frac{\partial L}{\partial a^l} = \frac{\partial L}{\partial a^{(l+1)}} \frac{\partial a^{(l+1)}}{\partial a^{(l)}}$$

- This is the basic principle behind backpropagation, we propagate the gradients **from the output layer back to the input layer**.

BACKPROPAGATION AND DERIVATIVES

BACKPROPAGATION ALGORITHM

- The backpropagation algorithm consists of the following steps:

STEP 1

Perform a forward pass to compute the activations and outputs of each layer $z^{(l)}$ and $a^{(l)}$

STEP 2

Compute the loss $L(y, \hat{y})$

STEP 3

Compute the gradient of the loss with respect to the output layer

$$\delta^{(L)} = \frac{\partial L}{\partial z^{(L)}}$$

BACKPROPAGATION AND DERIVATIVES

BACKPROPAGATION ALGORITHM

STEP 4

For each layer l from $L - 1$ to 1 :

- Compute the gradient of the loss with respect to the pre-activation z^l

$$\delta^{(l)} = (W^{(l+1)})^T \delta^{(l+1)} * \Phi'(z^{(l)})$$

- Compute the gradients with respect to the weights W_l and biases b_l

$$\frac{\partial L}{\partial W^{(l)}} = \delta^{(l)} (a^{(l-1)})^T$$

$$\frac{\partial L}{\partial b^{(l)}} = \delta^{(l)}$$

We can then use these gradients to update the parameters using gradient descent

BACKPROPAGATION AND DERIVATIVES

EXAMPLE: MLP WITH ONE HIDDEN LAYER

- Consider a MLP with one hidden layer and MSE loss function

STEP 1 - FORWARD PASS

$$z^{(1)} = W_1 x + b_1, \quad a^{(1)} = \Phi(z^{(1)})$$

$$z^{(2)} = W_2 a^{(1)} + b_2, \quad \hat{y} = a^{(2)} = \Phi(z^{(2)})$$

STEP 2

$$L(y, \hat{y}) = \frac{1}{2} \| y - \hat{y} \|^2$$

STEP 3 - OUTPUT LAYER GRADIENT

$$\delta^{(2)} = \frac{\partial L}{\partial z^{(2)}} = (\hat{y} - y) * \Phi'(z^{(2)})$$

STEP 4 - HIDDEN LAYER GRADIENT

$$\delta^{(1)} = (W_2)^T \delta^{(2)} * \Phi'(z^{(1)})$$

BACKPROPAGATION AND DERIVATIVES

EXAMPLE: MLP WITH ONE HIDDEN LAYER

FINAL GRADIENTS

$$\frac{\partial L}{\partial W_2} = \delta^{(2)} (a^{(1)})^T, \quad \frac{\partial L}{\partial b_2} = \delta^{(2)}$$

$$\frac{\partial L}{\partial W_1} = \delta^{(1)} x^T, \quad \frac{\partial L}{\partial b_1} = \delta^{(1)}$$

PARAMETER UPDATE


$$W_l := W_l - \eta \frac{\partial L}{\partial W_l}$$

$$b_l := b_l - \eta \frac{\partial L}{\partial b_l}$$

BACKPROPAGATION AND DERIVATIVES

VANISHING AND EXPLODING GRADIENTS

- During backpropagation, gradients can sometimes become very small (vanishing gradients) or very large (exploding gradients)
- Vanishing gradients make it difficult for the network to learn, especially in deep networks, as the updates to the weights become negligible
- Exploding gradients can lead to unstable training, causing the weights to oscillate or diverge
- Techniques to mitigate these issues include careful weight initialization, using activation functions like ReLU, and employing normalization techniques such as batch normalization



OPTIMIZATION ALGORITHMS FOR TRAINING NEURAL NETWORKS

OPTIMIZATION ALGORITHM

MOTIVATIONS

- Now that we can compute the prediction of the neural network and the derivatives of the loss function with respect to the parameters of the network
- To optimize the parameters, we need to solve the following optimization problem:

$$\min_{\theta} L(y, \hat{y}(\theta))$$

- Due to the large number of parameters and the complexity of the loss landscape, efficient optimization algorithms are crucial for effective training

OPTIMIZATION ALGORITHM

GRADIENT DESCENT

- We can use standard gradient descent methods to optimize the parameters of the neural network
- The update rule for gradient descent is given by:

$$\theta := \theta - \eta \nabla_{\theta} L$$

where η is the learning rate, controlling the step size of the updates

- In practice, computing the gradient over the entire dataset can be computationally expensive, leading to the use of stochastic and mini-batch gradient descent
- The complete gradient over the full dataset is given by:

$$\nabla_{\theta} L = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} L(y_i, \hat{y}_i(\theta))$$

but can be costly to be computed for large datasets

OPTIMIZATION ALGORITHM

CHALLENGES IN OPTIMIZATION

- The loss landscape of neural networks is often highly non-convex, with many local minima, saddle points, and flat regions
- This complexity can make optimization challenging, as standard gradient descent methods may get stuck in poor local minima or take a long time to converge
- Additionally, the choice of learning rate is critical; too high a learning rate can lead to divergence, while too low a learning rate can result in slow convergence
- Various optimization algorithms have been developed to address these challenges, including stochastic gradient descent, mini-batch gradient descent, momentum-based methods, and adaptive learning rate methods like Adam

OPTIMIZATION ALGORITHM

STOCHASTIC GRADIENT DESCENT

- Stochastic Gradient Descent (SGD) updates the parameters using the gradient computed from a batch of training examples

$$\nabla_{\theta} L \approx \frac{1}{n} \sum_{i \in \text{batch}} \nabla_{\theta} L(y_i, \hat{y}_i(\theta))$$

- The update rule for SGD is given by:

$$\theta := \theta - \eta \nabla_{\theta} L(y_i, \hat{y}_i(\theta))$$

- Advantages
 - Introduces noise into the updates, which can help escape local minima and improve generalization
 - Less computationally expensive per iteration, allowing for faster updates
- Downsides
 - Can become unstable and oscillate around minima if the learning rate is not properly tuned
 - Can be slower to converge if there is too much noise (i.e., batch size is too small)

OPTIMIZATION ALGORITHM

MOMENTUM SGD

- Momentum SGD introduces a velocity term to the parameter updates, helping to smooth out the updates and accelerate convergence
- The update rule for Momentum SGD is given by:

$$v_t := \beta v_{t-1} + (1 - \beta) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))$$

$$\theta_{t+1} := \theta_t - \eta v_t$$

where β is the momentum coefficient (typically between 0.9 and 0.99)

- Advantages
 - Helps accelerate convergence in the relevant direction and dampens oscillations
 - Can lead to faster training times compared to standard SGD
- Downsides
 - Requires tuning of the momentum coefficient in addition to the learning rate

OPTIMIZATION ALGORITHM

NESTEROV ACCELERATION

- Nesterov Acceleration is a variant of momentum SGD that looks ahead at the future position of the parameters to compute the gradient
- The update rule for NAG is given by:

$$v_t := \beta v_{t-1} + (1 - \beta) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t - \eta \beta v_{t-1}))$$

$$\theta_{t+1} := \theta_t - \eta v_t$$

- Advantages
 - Provides a more accurate estimate of the gradient, leading to improved convergence rates
 - Can help avoid overshooting minima compared to standard momentum SGD
- Downsides
 - Slightly more complex to implement and requires tuning of the momentum coefficient

OPTIMIZATION ALGORITHM

ADA GRAD

- AdaGrad adapts the learning rate for each parameter based on the historical gradients, allowing for larger updates when required
- The update rule for AdaGrad is given by:

$$G_t := \sum_{\tau=1}^t \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))^2$$

$$\theta_{t+1} := \theta_t - \left(\frac{\eta}{\sqrt{G_t} + \varepsilon} \right) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))$$

- Advantages
 - Automatically adjusts the learning rate, reducing the need for manual tuning
 - Particularly effective for sparse data and features
- Downsides
 - The learning rate can become excessively small over time, leading to slow convergence

OPTIMIZATION ALGORITHM

RMSProp

- RMSProp modifies AdaGrad by introducing a decay factor to the accumulated squared gradients, preventing the learning rate from becoming too small
- The update rule for RMSProp is given by:

$$E[g^2]_t := \beta E[g^2]_{t-1} + (1 - \beta) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))^2$$

$$\theta_{t+1} := \theta_t - \left(\frac{\eta}{\sqrt{E[g^2]_t} + \varepsilon} \right) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))$$

- Advantages
 - Maintains a more stable learning rate over time compared to AdaGrad
 - Effective for non-stationary objectives and online learning scenarios
- Downsides
 - Requires tuning of the decay factor in addition to the learning rate

OPTIMIZATION ALGORITHM

ADAM OPTIMIZER

- Adam (Adaptive Moment Estimation) combines the benefits of Momentum SGD and RMSProp, maintaining both a running average of the gradients and their squared values
- The update rule for Adam is given by:

$$m_t := \beta_1 m_{t-1} + (1 - \beta_1) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))$$

$$v_t := \beta_2 v_{t-1} + (1 - \beta_2) \nabla_{\theta} L(y_i, \hat{y}_i(\theta_t))^2$$

$$\widehat{m}_t := \frac{m_t}{1 - \beta_1^t}$$

$$\widehat{v}_t := \frac{v_t}{1 - \beta_2^t}$$

$$\theta_{t+1} := \theta_t - \left(\frac{\eta}{\sqrt{\widehat{v}_t} + \varepsilon} \right) \widehat{m}_t$$

OPTIMIZATION ALGORITHM

ADAM OPTIMIZER

- Advantages
 - Combines the benefits of momentum and adaptive learning rates
 - Generally performs well across a wide range of problems with minimal hyperparameter tuning
- Downsides
 - More complex to implement and requires tuning of multiple hyperparameters

OPTIMIZATION ALGORITHM

SUMMARY OF OPTIMIZATION ALGORITHMS

Algorithm	Key idea	Advantages	Disadvantages
SGD	Mini-batch	Simple, better generalization	Sensitive to η
Momentum SGD	Adds velocity	Accelerates convergence, reduces oscillations	Can oscillate, requires tuning of β
Nesterov Acceleration	Looks ahead to compute gradients	Improved convergence rates, avoids overshooting	More complex implementation
AdaGrad	Adapts learning rates based on historical gradients	Automatic learning rate adjustment, effective for sparse data	Learning rate can become too small
RMSProp	Introduces decay factor to AdaGrad	Stable learning rates, effective for non-stationary objectives	Requires tuning of decay factor
Adam	Combines momentum and adaptive learning rates	Performs well across various problems	Complex implementation, multiple hyperparameters

CONCLUSION

CONCLUSION

SUMMARY OF KEY POINTS

- Multi-layer perceptrons (MLPs) are powerful models capable of learning complex patterns through layers of interconnected neurons
- Forward propagation allows us to compute the output of the network given an input

$$\hat{y} = f_{\theta}(x)$$

- Backpropagation efficiently computes gradients for training

$$\nabla_{\theta} L(\theta) = \text{backpropagation}(y, \hat{y}(\theta))$$

- Various optimization algorithms, such as SGD, Momentum, RMSProp, and Adam, help in effectively training neural networks by addressing challenges in the loss landscape

$$\min_{\theta} L(y, \hat{y}(\theta))$$

CONCLUSION

CURRENT CHALLENGES AND FUTURE DIRECTIONS

- Despite the success of MLPs, several challenges remain in training deep neural networks
 1. Vanishing/exploding gradients: Developing techniques to mitigate these issues in very deep networks
 2. Overfitting: Developing better regularization techniques to improve generalization
 3. Scalability: Designing algorithms that can efficiently handle larger datasets and models
 4. Interpretability: Enhancing the transparency of neural networks to understand their decision-making processes
- Future research directions include exploring new architectures, optimization techniques, and applications across various domains

CONCLUSION

NEXT PRACTICAL

- In the next class, we will explore practical implementations of the concepts covered today
- We will implement a multi-layer perceptron from scratch, including forward propagation, backpropagation, and optimization using different algorithms
- We will also apply these techniques to real-world datasets, such as digit recognition using the MNIST dataset

CONCLUSION

NEXT LECTURE

- In the next lecture, we will delve into different neural network architectures, focusing on Convolutional Neural Networks (CNNs)
- We will explore their architecture, how they differ from MLPs, and why they are particularly effective for image processing tasks
- We will also discuss advanced optimization techniques and regularization methods (Dropout, BatchNorm...)